

Contents

- Understanding Gradient Descent
 - A minimal implementation
 - Choosing the learning rate
 - Variants worth knowing

Understanding Gradient Descent

Great question! Here's a clear walkthrough of how gradient descent actually works .

The core idea: we want to minimise a loss function $L(\theta)$ by repeatedly stepping in the direction that reduces it fastest - the negative gradient. Each update looks like:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} L(\theta_t)$$

where η is the **learning rate**.

A minimal implementation

```
import numpy as np

def gradient_descent(grad, theta, lr=0.1, steps=100):
    for _ in range(steps):
        theta = theta - lr * grad(theta)
    return theta
```

Choosing the learning rate

- Too **small** → training crawls and may never converge in reasonable time
- Too **large** → the loss oscillates or diverges entirely
- A good default is to start around $1e-3$ and tune from there

Variants worth knowing

Variant	What changes	Good for
SGD	one sample per step	large datasets
Momentum	accumulates past gradients	ravines / plateaus
Adam	per-parameter adaptive rates	most deep-learning defaults